

1. **Implement a simple neural network using tensorflow to classify the handwritten digits from MNIST dataset**



```
1 # Step 1: Import Libraries
2 import tensorflow as tf # TensorFlow for machine learning
3 from tensorflow.keras import layers, models # Import Keras layers and model APIs
4 import matplotlib.pyplot as plt # For visualizing predictions
5
6 # Step 2: Load and Preprocess MNIST Dataset
7 mnist = tf.keras.datasets.mnist # Load the built-in MNIST dataset
8 (x_train, y_train), (x_test, y_test) = mnist.load_data() # Split into training and test sets
9
10 # Normalize pixel values to [0, 1] for faster training and better performance
11 x_train, x_test = x_train / 255.0, x_test / 255.0
12
13 # Step 3: Build the Neural Network Model
14 model = models.Sequential([ # Create a sequential model
15     layers.Flatten(input_shape=(28, 28)), # Flatten 28x28 images to 784-element vectors
16     layers.Dense(128, activation='relu'), # Fully connected hidden layer with 128 units and ReLU activation
17     layers.Dropout(0.2), # Dropout layer to prevent overfitting (drops 20% of nodes randomly)
18     layers.Dense(10, activation='softmax') # Output layer: 10 neurons (one per digit), softmax for probabilities
19 ])
20
21 # Step 4: Compile the Model
22 model.compile(
23     optimizer='adam', # Use Adam optimizer (adaptive learning rate)
24     loss='sparse_categorical_crossentropy', # Suitable for integer labels (0-9)
25     metrics=['accuracy'] # Track accuracy during training and testing
26 )
27
28 # Step 5: Train the Model
29 model.fit(
30     x_train, y_train, # Input and labels
31     epochs=5, # Number of training passes over the data
32     validation_split=0.1 # 10% of training data is used for validation
33 )
34
35 # Step 6: Evaluate on Test Data
36 test_loss, test_acc = model.evaluate(x_test, y_test) # Evaluate on unseen test data
37 print('\nTest accuracy:', test_acc) # Print the final test accuracy
38
39 # Optional: Step 7 - Plot Predictions
40 predictions = model.predict(x_test) # Predict class probabilities for each test image
41
42 # Define a function to display a prediction
43 def plot_image(i, predictions_array, true_label, img):
44     plt.grid(False) # Remove grid lines
45     plt.xticks([]) # Remove x-axis ticks
46     plt.yticks([]) # Remove y-axis ticks
47     plt.imshow(img, cmap=plt.cm.binary) # Show image in grayscale
48
49     predicted_label = tf.argmax(predictions_array) # Get predicted digit (highest probability)
50     color = 'blue' if predicted_label == true_label else 'red' # Blue if correct, red if wrong
51
52     plt.xlabel(f"Pred: {predicted_label} (True: {true_label})", color=color) # Add label with prediction
53
54 # Display the first 10 test images with predictions
55 plt.figure(figsize=(10, 5)) # Set figure size
56 for i in range(10):
57     plt.subplot(2, 5, i + 1) # 2 rows x 5 columns of subplots
58     plot_image(i, predictions[i], y_test[i], x_test[i]) # Show image and prediction
59 plt.tight_layout() # Adjust layout to avoid overlap
60 plt.show() # Display the full plot
61
```

2. Train a convolutional neural network to classify dogs and cats using kaggle dataset

```
1  # Step 1: Import required libraries
2  import tensorflow as tf
3  from tensorflow.keras.preprocessing.image import ImageDataGenerator
4
5  # Step 2: Preprocess the images (rescale pixel values to [0, 1])
6  train_gen = ImageDataGenerator(rescale=1./255)
7  val_gen = ImageDataGenerator(rescale=1./255)
8
9  # Step 3: Load the training and validation images
10 train_data = train_gen.flow_from_directory(
11     'data/train',          # Folder should contain 'cats' and 'dogs' subfolders
12     target_size=(150, 150), # Resize all images to 150x150
13     batch_size=20,         # Process 20 images at a time
14     class_mode='binary'    # 0 = cat, 1 = dog
15 )
16
17 val_data = val_gen.flow_from_directory(
18     'data/validation',     # Validation folder with same structure
19     target_size=(150, 150),
20     batch_size=20,
21     class_mode='binary'
22 )
23
24 # Step 4: Build the CNN model
25 model = tf.keras.models.Sequential([
26     tf.keras.layers.Conv2D(32, (3, 3), activation='relu', input_shape=(150, 150, 3)),
27     tf.keras.layers.MaxPooling2D(2, 2),
28
29     tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
30     tf.keras.layers.MaxPooling2D(2, 2),
31
32     tf.keras.layers.Flatten(),
33     tf.keras.layers.Dense(128, activation='relu'),
34     tf.keras.layers.Dense(1, activation='sigmoid') # Binary output
35 ])
36
37 # Step 5: Compile the model
38 model.compile(
39     loss='binary_crossentropy',
40     optimizer='adam',
41     metrics=['accuracy']
42 )
43
44 # Step 6: Train the model
45 model.fit(
46     train_data,
47     epochs=5,
48     validation_data=val_data
49 )
50
```

3.fine tune a pre trained CNN model Model V66 RESNET or mobile net on small dataset for a specific classification tasks

```

1  # Import required libraries
2  import tensorflow as tf
3  from tensorflow.keras.preprocessing.image import ImageDataGenerator
4
5  # Step 1: Preprocess image data
6  # Rescale pixel values to [0, 1]
7  train_gen = ImageDataGenerator(rescale=1./255)
8  val_gen = ImageDataGenerator(rescale=1./255)
9
10 # Load training images from folders
11 train_data = train_gen.flow_from_directory(
12     'data/train',          # Folder with class subfolders (e.g., cats, dogs)
13     target_size=(160, 160), # Resize images
14     batch_size=32,
15     class_mode='binary'    # For 2-class classification
16 )
17
18 # Load validation images
19 val_data = val_gen.flow_from_directory(
20     'data/validation',
21     target_size=(160, 160),
22     batch_size=32,
23     class_mode='binary'
24 )
25
26 # Step 2: Load pre-trained MobileNetV2 (without top layer)
27 base_model = tf.keras.applications.MobileNetV2(
28     input_shape=(160, 160, 3),
29     include_top=False,
30     weights='imagenet'
31 )
32 base_model.trainable = False # Freeze the base model layers
33
34 # Step 3: Add custom classification layers on top
35 model = tf.keras.Sequential([
36     base_model,
37     tf.keras.layers.GlobalAveragePooling2D(), # Reduces tensor shape
38     tf.keras.layers.Dense(1, activation='sigmoid') # Output: 0 or 1
39 ])
40
41 # Step 4: Compile the model
42 model.compile(
43     optimizer='adam',
44     loss='binary_crossentropy',
45     metrics=['accuracy']
46 )
47
48 # Step 5: Train top layers (feature extractor phase)
49 model.fit(
50     train_data,
51     epochs=5,
52     validation_data=val_data
53 )
54
55 # Step 6: Unfreeze some base model layers for fine-tuning
56 base_model.trainable = True
57 fine_tune_at = 100 # Freeze all layers before this
58
59 for layer in base_model.layers[:fine_tune_at]:
60     layer.trainable = False
61
62 # Step 7: Recompile with a lower learning rate
63 model.compile(
64     optimizer=tf.keras.optimizers.Adam(learning_rate=1e-5),
65     loss='binary_crossentropy',
66     metrics=['accuracy']
67 )
68
69 # Step 8: Continue training (fine-tuning phase)
70 model.fit(
71     train_data,
72     epochs=10,
73     initial_epoch=5,

```

```
73         validation_data=val_data,
74         validation_data=val_data
75     )
76
```

1. Create RNN or a transformer to classify text documents into different categories such as sentiment analysis or topic selection

```
1  # Import required libraries
2  import tensorflow as tf
3  from tensorflow.keras import layers
4  from tensorflow.keras.datasets import imdb
5  from tensorflow.keras.preprocessing.sequence import pad_sequences
6
7  # Step 1: Load IMDB sentiment dataset
8  # Keep only the 10,000 most frequent words
9  vocab_size = 10000
10 max_length = 200
11
12 (x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=vocab_size)
13
14 # Step 2: Pad all sequences to the same length
15 x_train = pad_sequences(x_train, maxlen=max_length)
16 x_test = pad_sequences(x_test, maxlen=max_length)
17
18 # Step 3: Build the RNN model
19 model = tf.keras.Sequential([
20     layers.Embedding(input_dim=vocab_size, output_dim=32, input_length=max_length), # Word embedding
21     layers.SimpleRNN(32), # RNN layer
22     layers.Dense(1, activation='sigmoid') # Output: 0 (neg) or 1 (pos)
23 ])
24
25 # Step 4: Compile the model
26 model.compile(optimizer='adam',
27               loss='binary_crossentropy',
28               metrics=['accuracy'])
29
30 # Step 5: Train the model
31 model.fit(x_train, y_train, epochs=5, validation_split=0.2)
32
33 # Step 6: Evaluate on test data
34 loss, accuracy = model.evaluate(x_test, y_test)
35 print("Test Accuracy:", accuracy)
36
```

5. Generate Synthetic images or data that resemble real data using GAN


```

1  # Simple GAN to Generate Handwritten Digits
2
3  import tensorflow as tf
4  from tensorflow.keras import layers
5  import numpy as np
6  import matplotlib.pyplot as plt
7
8  # Load and preprocess MNIST data
9  (x_train, _), _ = tf.keras.datasets.mnist.load_data()
10 x_train = (x_train - 127.5) / 127.5 # Normalize to [-1, 1]
11 x_train = x_train.reshape(-1, 28, 28, 1).astype('float32')
12 train_ds = tf.data.Dataset.from_tensor_slices(x_train).shuffle(60000).batch(256)
13
14 # Generator model: turns noise into an image
15 def build_generator():
16     model = tf.keras.Sequential([
17         layers.Dense(7*7*256, use_bias=False, input_shape=(100,)),
18         layers.BatchNormalization(), layers.LeakyReLU(),
19         layers.Reshape((7, 7, 256)),
20         layers.Conv2DTranspose(128, 5, strides=1, padding='same', use_bias=False),
21         layers.BatchNormalization(), layers.LeakyReLU(),
22         layers.Conv2DTranspose(64, 5, strides=2, padding='same', use_bias=False),
23         layers.BatchNormalization(), layers.LeakyReLU(),
24         layers.Conv2DTranspose(1, 5, strides=2, padding='same', use_bias=False, activation='tanh')
25     ])
26     return model
27
28 # Discriminator model: tells real from fake images
29 def build_discriminator():
30     model = tf.keras.Sequential([
31         layers.Conv2D(64, 5, strides=2, padding='same', input_shape=[28,28,1]),
32         layers.LeakyReLU(), layers.Dropout(0.3),
33         layers.Conv2D(128, 5, strides=2, padding='same'),
34         layers.LeakyReLU(), layers.Dropout(0.3),
35         layers.Flatten(), layers.Dense(1)
36     ])
37     return model
38
39 # Losses and optimizers
40 cross_entropy = tf.keras.losses.BinaryCrossentropy(from_logits=True)
41 generator = build_generator()
42 discriminator = build_discriminator()
43 gen_opt = tf.keras.optimizers.Adam(1e-4)
44 disc_opt = tf.keras.optimizers.Adam(1e-4)
45
46 # Training step
47 @tf.function
48 def train_step(images):
49     noise = tf.random.normal([256, 100])
50     with tf.GradientTape() as gen_tape, tf.GradientTape() as disc_tape:
51         gen_imgs = generator(noise, training=True)
52         real_out = discriminator(images, training=True)
53         fake_out = discriminator(gen_imgs, training=True)
54         gen_loss = cross_entropy(tf.ones_like(fake_out), fake_out)
55         disc_loss = cross_entropy(tf.ones_like(real_out), real_out) + \
56                     cross_entropy(tf.zeros_like(fake_out), fake_out)
57     gen_grad = gen_tape.gradient(gen_loss, generator.trainable_variables)
58     disc_grad = disc_tape.gradient(disc_loss, discriminator.trainable_variables)

```

```

59     gen_opt.apply_gradients(zip(gen_grad, generator.trainable_variables))
60     disc_opt.apply_gradients(zip(disc_grad, discriminator.trainable_variables))
61
62     # Generate and show images
63     def generate_images(model, test_input):
64         preds = model(test_input, training=False)
65         plt.figure(figsize=(4,4))
66         for i in range(preds.shape[0]):
67             plt.subplot(4, 4, i+1)
68             plt.imshow((preds[i, :, :, 0] + 1) / 2.0, cmap='gray')
69             plt.axis('off')
70         plt.tight_layout()
71         plt.show()
72
73     # Train loop
74     def train(dataset, epochs):
75         seed = tf.random.normal([16, 100])
76         for epoch in range(epochs):
77             for batch in dataset:
78                 train_step(batch)
79                 print(f'Epoch {epoch+1} done')
80                 generate_images(generator, seed)
81
82     # Run training
83     train(train_ds, epochs=3) # You can increase epochs for better results
84

```

6.implementation object detection model using frameworks like tensorflow or pytorch and evaluate it performance on dataset such as cow or pascal



```
1 # Step 1: Install required packages (run this in your terminal or Jupyter notebook)
2 # !pip install tensorflow tensorflow-hub opencv-python
3
4 # Step 2: Import libraries
5 import tensorflow as tf
6 import tensorflow_hub as hub
7 import numpy as np
8 import cv2
9 import matplotlib.pyplot as plt
10
11 # Step 3: Load and preprocess one image
12 def load_image(img_path):
13     img = tf.io.read_file(img_path)
14     img = tf.image.decode_jpeg(img, channels=3)
15     img = tf.image.resize(img, (384, 384)) / 255.0 # Resize and normalize
16     return img
17
18 img_path = 'cow.jpg' # Replace with your image file
19 image = load_image(img_path)
20 input_tensor = tf.expand_dims(image, 0) # Add batch dimension
21
22 # Step 4: Load pre-trained object detection model
23 model = hub.load("https://tfhub.dev/tensorflow/efficientdet/lite2/detection/1")
24
25 # Step 5: Run detection
26 output = model(input_tensor)
27 boxes = output["detection_boxes"][0].numpy()
28 classes = output["detection_classes"][0].numpy().astype(np.int32)
29 scores = output["detection_scores"][0].numpy()
30
31 # Step 6: Draw boxes on the image
32 labels = {1: 'person', 17: 'cat', 18: 'dog', 20: 'cow'} # Add more as needed
33
34 img_np = np.array(image * 255, dtype=np.uint8)
35 h, w = img_np.shape[:2]
36
37 for box, cls, score in zip(boxes, classes, scores):
38     if score < 0.3:
39         continue
40     y1, x1, y2, x2 = box
41     start = (int(x1 * w), int(y1 * h))
42     end = (int(x2 * w), int(y2 * h))
43     cv2.rectangle(img_np, start, end, (0, 255, 0), 2)
44     label = f"{labels.get(cls, 'ID:'+str(cls))} ({score:.2f})"
45     cv2.putText(img_np, label, (start[0], start[1]-10), cv2.FONT_HERSHEY_SIMPLEX, 0.5, (255, 0, 0), 1)
46
47 # Step 7: Show result
48 plt.imshow(img_np)
49 plt.axis("off")
50 plt.show()
51
```